

TRINIDAD BORRELL

✉ triniborrell@gmail.com

in linkedin.com/in/trinidad-borrell

github github.com/trinidadBorrell

Summary

I'm a physicist turned neuroscientist and deep learning researcher, currently a PhD candidate at the Paris Brain Institute and Forschungszentrum Jülich. My trajectory runs from the mathematical foundations of particle physics, through the study of human and machine robustness to adversarial attacks, and into deep learning for clinical brain signals. This has meant working hands-on with transformer-based foundation models and graph variational autoencoders on EEG data from patients with disorders of consciousness. Additionally, I am interested in AI safety, particularly mechanistic interpretability, and was selected as a MARS V Fellow, where I lead a project on characterizing the persona space of large language models.

Experience

AI Safety Fellow

Jun 2026 – Today

MARS V Programme

Oxford, England

- Selected for a highly competitive AI safety research mentorship programme (5% acceptance rate) run by the Cambridge AI Safety Hub, collaborating in a team to produce publishable work. Conducting mechanistic interpretability research on persona manifolds in large language models. Mentored by Thomas Jiralerspong.

Doctoral Candidate

Oct 2025 – Today

Paris Brain Institute | Forschungszentrum Jülich

Paris, France | Jülich, Germany

- Benchmarked seven time-series and EEG foundation models (TOTEM, LaBram, NeuroLM, CbraMod, EEGPT, BIOT and SleepFM) for classification of patients with disorders of consciousness, establishing the first systematic evaluation framework for this clinical population; Decoder-only based and EEG pretrained models consistently outperformed general-purpose time-series models across the majority of metrics. github.com/doc-benchmark
- Designed and trained a Graph Autoencoder over EEG functional connectivity matrices using GraphSAGE based message passing and subject-level data splitting, demonstrating that the latent space captures meaningful brain connectivity patterns across patient states. github.com/gnn-connectivity
- Built a hybrid CNN and GNN pipeline for context-aware cancer cell classification on the NuCLS histopathology dataset, encoding spatial relationships between nuclei as graph edges and developing a class-balanced data augmentation strategy to mitigate severe label imbalance across 8 cell types; achieved 0.65 AUC, competitive with pretrained CNN baselines github.com/gnn-cancer-cell

AI Engineer

Apr 2025 – Oct 2025

Paris Brain Institute

Paris, France

- Developed preprocessing and marker computation pipelines for multi-subject clinical EEG data. Containerized the full pipeline with Docker to enable reproducible execution across research environments.

AI Researcher

Feb 2024 – Apr 2025

Consciousness, Culture and Complexity lab

Buenos Aires, Argentina

- Investigated the differential impact of adversarial perturbations on human visual perception and CNNs (EfficientNet, ResNet, Inception V3); presented as first-author at Khipu Latin American Conference in AI 2025 github.com/adv-attacks
- Developed a LLM and RAG pipeline over transcribed mind-wandering patient recordings, enabling structured extraction and analysis of spontaneous thought patterns from clinical interview data.

Research Intern

Mar 2023 – Mar 2024

Lambda Lab

Buenos Aires, Argentina

- Assessed coded aperture techniques for X-ray detection through Monte Carlo-based computational analysis github.com/coded-aperture

Education

Sorbonne Université

Sept 2025 – Today

PhD Candidate in Neuroscience

Paris, France

University of Buenos Aires

Mar 2019 – Mar 2025

Licentiate in Physics (Equivalent to Bachelor and Master on Science)

Buenos Aires, Argentina

GPA (in 10-point scale): 9.4

Collaborations

Platonic Universe

Nov 2025 – Today

Universe TBD

- Investigating whether foundation models converge toward a shared representational structure as scale increases, using astronomical data as input.

Teaching & Mentoring Experience

AI Safety Mentor

Aug 2026 – Today

SPAR

Remote

- Mentoring SPAR fellows on whether curvature-aware methods reveal that concept geometry is shared across different language models.

Technical AI Safety Facilitator

May 2026 – Today

BlueDot

Remote

- Guiding cohorts of emerging AI safety researchers through core concepts in alignment, mechanistic interpretability, evaluations, and AI control.

Data Science Teacher

Jul 2021 – May 2022

CoderHouse

Buenos Aires, Argentina

- Delivered live online instruction in data science, teaching Python, data analysis, visualization, and machine learning through practical, project-based learning.

Papers

T. Borrell, S. Dillmann, K. Duraphe, F. Eris, A. Khederlarian, A. Kumar, G. Marraffini, M.J. Smith, S. Sourav, R. Di Tella, and JF. Wu. The Platonic Universe: Do Foundation Models See the Same Sky?. ICML 2026 Mech Interp Workshop.

T. Borrell, L. Belloli, G. Marraffini, K R Patil, J D. Sitt, and F. Raimondo. What Have EEG Foundation Models Actually Learned About the Conscious Brain?. *Submitted to Neurips 2026*

G. Marraffini, G. Mahuas, **T. Borrell**, V. Shevchenko, D. Wassermann. The Variance Brain Foundation Models Forgot: Third-Order Statistics Predict Cognition Where Billion-Parameter Models Fail. *Submitted to Neurips 2026*. arXiv:2606.04010.

M. Carro, D. Mester, F. Nieto, O. Stanchi, G. Bergman, G. Marraffini, **T. Borrell** et al. Humans Can Correct Their Judgements Through Debate, but Weaker AI Models Can Not. *Submitted to Neurips 2026*

Grants

BlueDot Rapid Grant. ICML 2026 Full Travel Grant.

Iliad Intensive 2026. Full Travel Grant.

BlueDot Rapid Grant. Compute Grant.

ML4Good Bootcamp 2026. Full Travel Grant and Scholarship.

Winter School AI for Sciences 2026 Full Travel Grant and Scholarship.

Paris Brain Institute America Doctoral Grant. 3-year PhD Grant.

Khipu Latin American Meeting in Artificial Intelligence 2025 Full Travel Grant.

Brain, Minds and Machines 2024 Summer School Full Travel Grant and Scholarship.

Presentations and Posters

T. Borrell, L. Belloli, G. Marraffini, K R Patil, J D. Sitt, and F. Raimondo. (June, 2026). *A Comprehensive Evaluation Framework for EEG Foundation Models in Disorders of Consciousness*. Poster presented at OHBM 2026. Bordeaux, France.

T. Borrell, E. Tagliazucchi. (March, 2025). *Subtle or obvious? Confidence reports indicate the success of adversarial attacks on humans*. Poster presented at Khipu Latin American Meeting in Artificial Intelligence 2025. Santiago, Chile.

T. Borrell, E. Tagliazucchi. (November, 2024). *Human and machine vulnerability to adversarial perturbations*. Poster presented at the Scientific Symposium on AI and Applications, led by the University of Saint Andrews. Buenos Aires, Argentina.

T. Borrell, E. Tagliazucchi. (September, 2024). *Human and machine vulnerability to adversarial perturbations*. Poster presented at the 2024 Annual Reunion of the Argentinian Physics Association. San Luis, Argentina.

T. Borrell, E. Sprejer, A. G. Magnoni, D. Rodrigues, Q. Pears Stefano. (September, 2023). *Impossible Photographs*. Poster presented at the 2023 Annual Reunion of the Argentinian Physics Asociation. Buenos Aires, Argentina.

Skills

AI/ML: Python, PyTorch, HuggingFace, FastAPI, Scikit-learn, Claude Code, Time Series Foundation Models, Fine-tuning, Evaluation, Optimization, LLMs, Transformers, RAG, LangGraph, LangChain Prompt Engineering, CNNs, Synthetic Data Generation, Agentic Engineering, GenAI.

Data: Pandas, NumPy, SQL, MongoDB.

Infrastructure: Docker, Git, Github, Docker, HTCondor.

Visualization: Plotly, Seaborn, Matplotlib.

Languages: Spanish (Native), English (TOEFL iBT: 106), French (DELF B1: 78).